

**ASSIGNMENT -- 7****NAME -- SHRASHTI YADAV****REG NO -- 20233269****SECTION -- D1**

```
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers, regularizers
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split

# Load MNIST dataset
(x_train, y_train), (x_test, y_test) = keras.datasets.mnist.load_data()

# Normalize pixel values to [0, 1]
x_train, x_test = x_train / 255.0, x_test / 255.0

# Flatten images into 1D arrays (28x28 = 784)
x_train = x_train.reshape(-1, 784)
x_test = x_test.reshape(-1, 784)

# One-hot encode labels
y_train = keras.utils.to_categorical(y_train, 10)
y_test = keras.utils.to_categorical(y_test, 10)

# Split training data (80% training, 20% validation)
x_train, x_val, y_train, y_val = train_test_split(x_train, y_train, test_size=0.2, random_state=42)

# Build ANN model with L2 regularization and Dropout
model = keras.Sequential([
    layers.Dense(256, activation='relu', kernel_regularizer=regularizers.l2(0.001), input_shape=(784,)),
    layers.Dropout(0.3),
    layers.Dense(128, activation='relu', kernel_regularizer=regularizers.l2(0.001)),
    layers.Dropout(0.3),
    layers.Dense(10, activation='softmax')
])

# Compile the model
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

# Train the model
history = model.fit(x_train, y_train, epochs=10, batch_size=32, validation_data=(x_val, y_val))

# Evaluate on test set
test_loss, test_acc = model.evaluate(x_test, y_test, verbose=2)
print(f"Test accuracy: {test_acc:.4f}")

# Plot training history
plt.figure(figsize=(12, 4))
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.title('Training vs Validation Accuracy')

plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.title('Training vs Validation Loss')

plt.show()
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.11490434/11490434> 0s 0us/step

```
/usr/local/lib/python3.11/dist-packages/keras/src/layers/core/dense.py:87: UserWarning:
  super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

Epoch 1/10  
 1500/1500 ————— 16s 9ms/step - accuracy: 0.8239 - loss: 0.8919 - val\_accu  
 Epoch 2/10  
 1500/1500 ————— 16s 6ms/step - accuracy: 0.9404 - loss: 0.3974 - val\_accu  
 Epoch 3/10  
 1500/1500 ————— 9s 6ms/step - accuracy: 0.9495 - loss: 0.3365 - val\_accu  
 Epoch 4/10  
 1500/1500 ————— 10s 6ms/step - accuracy: 0.9516 - loss: 0.3156 - val\_accu  
 Epoch 5/10  
 1500/1500 ————— 10s 6ms/step - accuracy: 0.9515 - loss: 0.3078 - val\_accu  
 Epoch 6/10  
 1500/1500 ————— 9s 6ms/step - accuracy: 0.9550 - loss: 0.2953 - val\_accu  
 Epoch 7/10  
 1500/1500 ————— 9s 6ms/step - accuracy: 0.9560 - loss: 0.2946 - val\_accu  
 Epoch 8/10  
 1500/1500 ————— 11s 6ms/step - accuracy: 0.9556 - loss: 0.2927 - val\_accu  
 Epoch 9/10  
 1500/1500 ————— 10s 6ms/step - accuracy: 0.9543 - loss: 0.2981 - val\_accu  
 Epoch 10/10  
 1500/1500 ————— 10s 6ms/step - accuracy: 0.9560 - loss: 0.2894 - val\_accu  
 313/313 - 1s - 3ms/step - accuracy: 0.9675 - loss: 0.2500  
 Test accuracy: 0.9675

